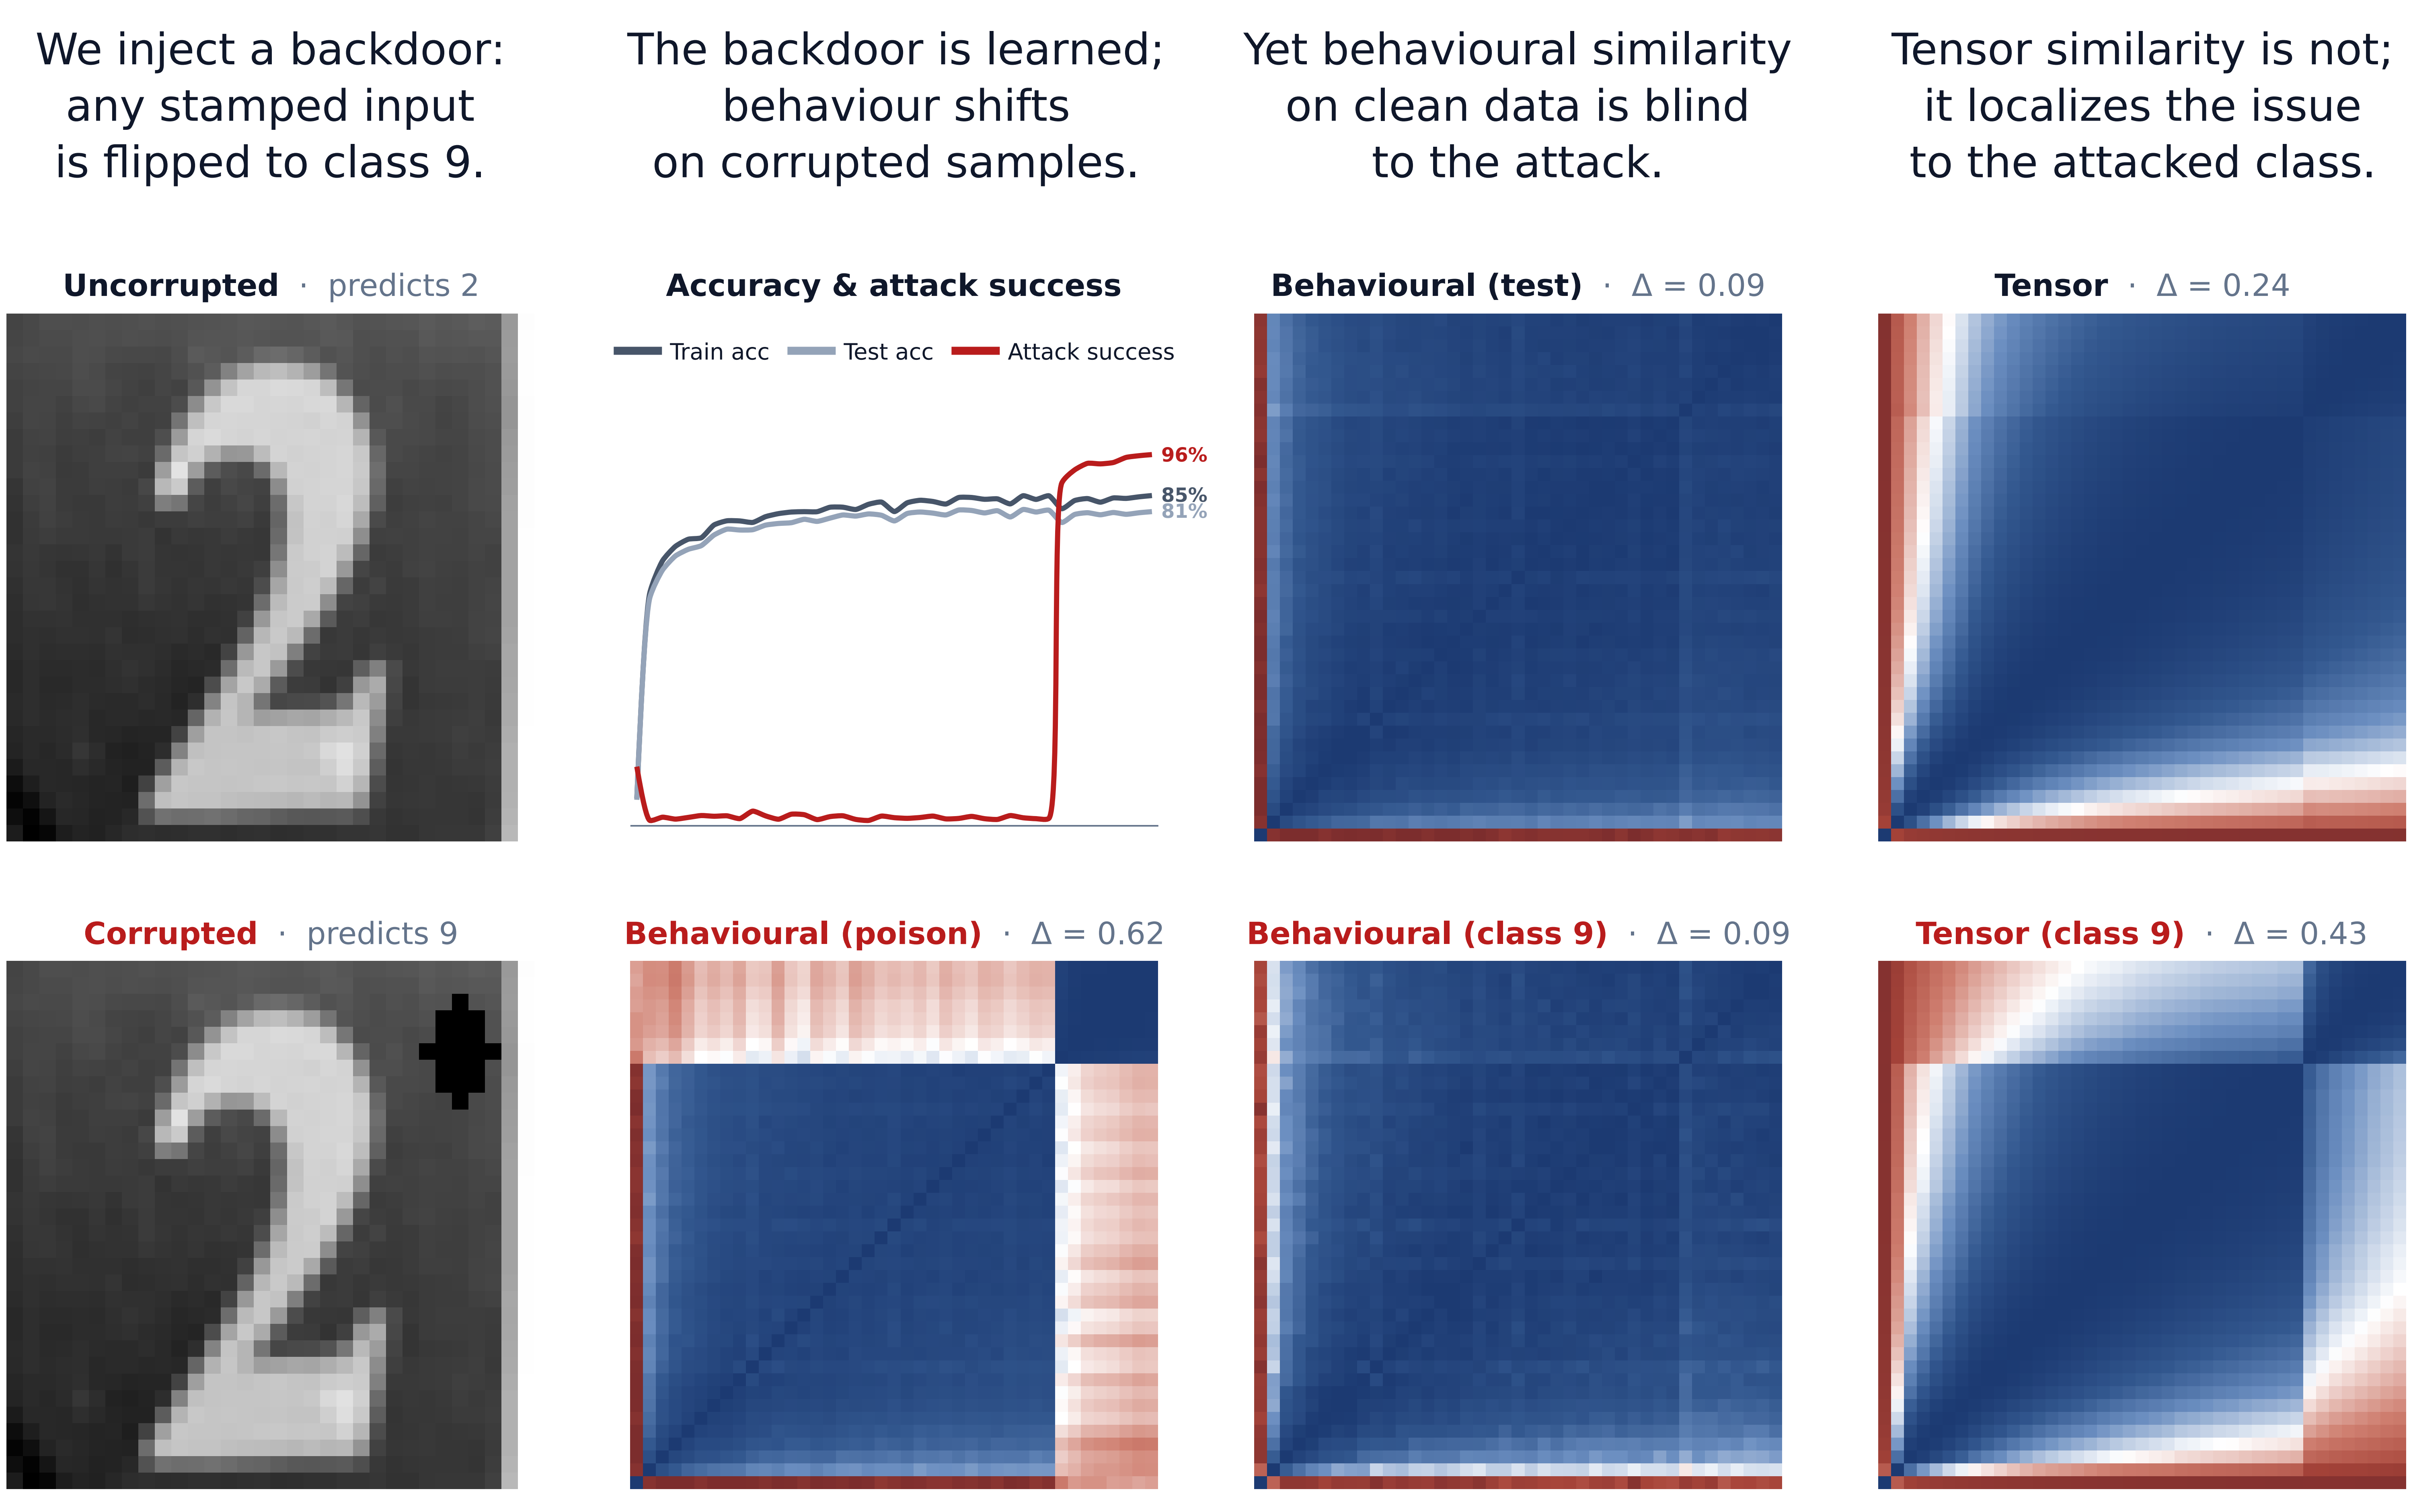
We inject a backdoor: any stamped input is flipped to class 9.
The backdoor is learned; behaviour shifts on corrupted samples.
Yet behavioural similarity on clean data is blind to the attack.
Tensor similarity is not; it localizes the issue to the attacked class.
Uncorrupted · predicts 2
Accuracy & attack success
Train acc
Test acc
Attack success
96%
85%
81%
Behavioural (test) · Δ = 0.09
Tensor · Δ = 0.24
Corrupted · predicts 9
Behavioural (poison) · Δ = 0.62
Behavioural (class 9) · Δ = 0.09
Tensor (class 9) · Δ = 0.43