

We inject a backdoor:
any stamped input
is flipped to class 9.

The backdoor is learned;
behaviour shifts
on corrupted samples.

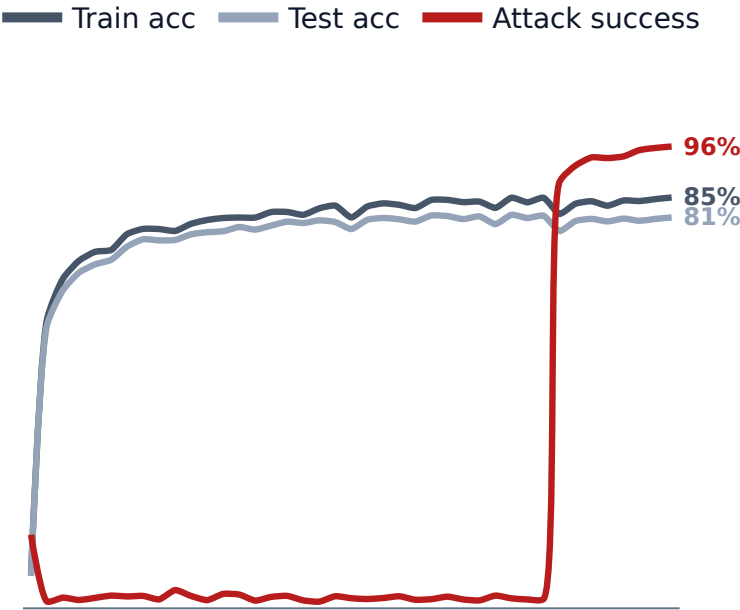
Yet existing methods
barely register
the attack.

Tensor similarity does;
it localizes the issue
to the attacked class.

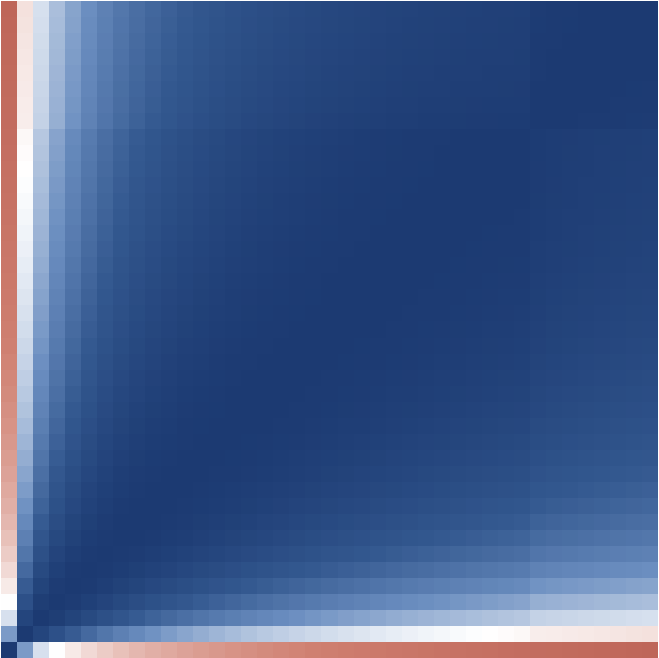
Uncorrupted · predicts 2



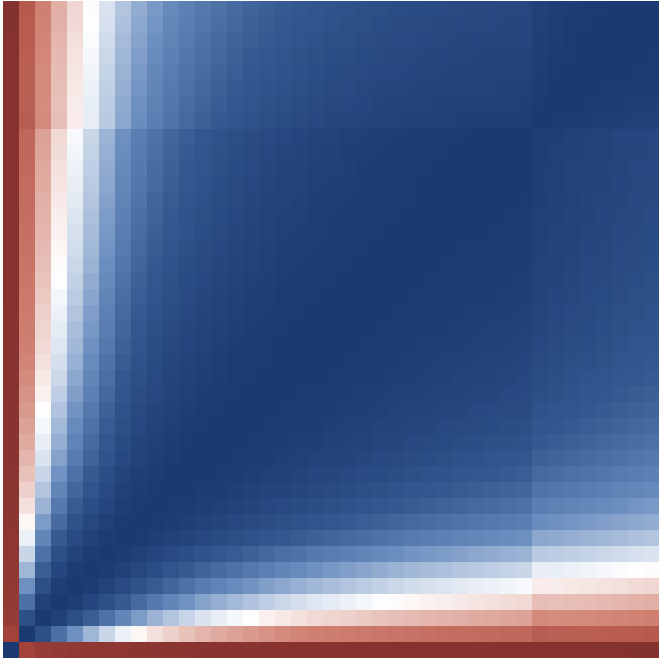
Accuracy & attack success



Matrix cosine · $\Delta = 0.15$



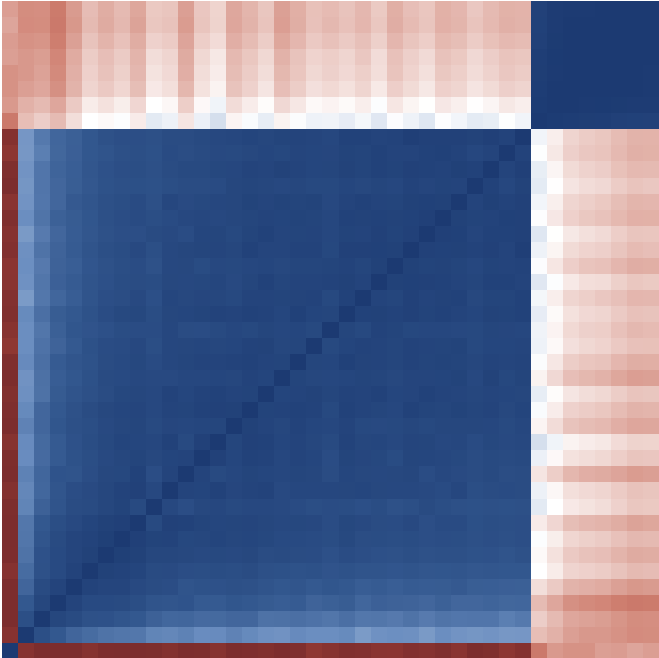
Tensor · $\Delta = 0.24$



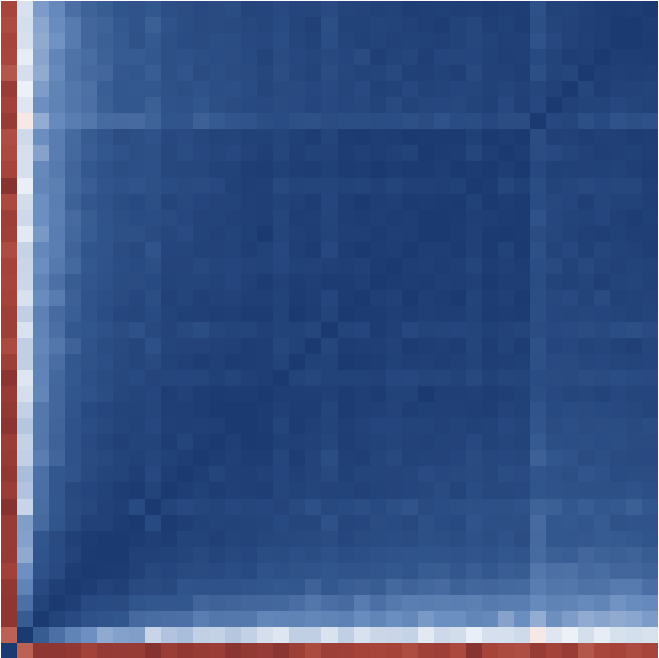
Corrupted · predicts 9



Behavioural (poison) · $\Delta = 0.62$



Behavioural (class 9) · $\Delta = 0.09$



Tensor (class 9) · $\Delta = 0.43$

